
SQL and NoSQL Tutorial

Design Lab
Jan 15, 2026

SQL: A Brief Introduction




Standard SQL Queries

- CREATE
- INSERT
- SELECT
- DELETE
- UPDATE

JOINS in SQL

- INNER JOIN: The INNER JOIN keyword returns only rows with a match in both tables.


```
SELECT column_name(s)
FROM table1
INNER JOIN table2
ON table1.column_name = table2.column_name;
```



JOINS in SQL

- LEFT JOIN: The LEFT JOIN keyword returns all records from the left table (table1), and the matching records from the right table (table2).


```
SELECT column_name(s)  
FROM table1  
LEFT JOIN table2  
ON table1.column_name = table2.column_name;
```



JOINS in SQL

- RIGHT JOIN: The RIGHT JOIN keyword returns all records from the right table, even if there are no matches in the left table.

```
SELECT column_name(s)
FROM table1
RIGHT JOIN table2
ON table1.column_name = table2.column_name;
```



Group By and HAVING clause

```
SELECT column_name(s)  
FROM table_name  
WHERE condition  
GROUP BY column_name(s)  
HAVING condition  
ORDER BY column_name(s);
```

ANY/ALL Operators

ANY operator:

- returns TRUE if ANY of the subquery values meet the condition

```
SELECT column_name(s)
FROM table_name
WHERE column_name operator ANY
  (SELECT column_name
   FROM table_name
   WHERE condition);
```

ALL operator:

- returns TRUE if ALL of the subquery values meet the condition

```
SELECT column_name(s)
FROM table_name
WHERE column_name operator ALL
  (SELECT column_name
   FROM table_name
   WHERE condition);
```

CREATE VIEW

- In SQL, a view is a virtual table based on the result-set of an SQL statement.
- A view contains rows and columns, just like a real table.
- You can add SQL statements and functions to a view and present the data as if the data were coming from one single table.

```
CREATE VIEW view_name AS
    SELECT column1, column2, ...
    FROM table_name
    WHERE condition;
```

Commit and Rollback

SQL Commit

COMMIT is the SQL command that is used for storing changes performed by a transaction.

SQL RollBack

ROLLBACK is the SQL command that is used for reverting changes performed by a transaction.

PARTITION BY in SQL

What is PARTITION BY?

- Divides rows into groups
- Used with window functions
- Calculation resets for each group
- Does not reduce rows

Syntax

```
FUNCTION() OVER (  
    PARTITION BY column  
    ORDER BY column  
)
```

PARTITION BY splits data into groups without collapsing rows

ROW_NUMBER() in SQL

What is ROW_NUMBER()?

- Assigns a **serial number** to each row
- Numbers start from 1
- Order is decided by **ORDER BY**
- **Does not remove rows**

Syntax

```
ROW_NUMBER() OVER
(
    PARTITION BY
    columns
    ORDER BY column
)
```

ROW_NUMBER() + filtering = controlled row selection

LEAD() in SQL

What is LEAD()?

- Accesses the **next row's value**
- No self-join needed
- Works within an **ordered set**

Ordered set = rows arranged in a specific order before the function runs.

Syntax

LEAD(column) OVER (ORDER BY column)

LEAD() looks ahead without joins

RANK() in SQL

What is RANK()?

- Assigns **rank** to rows
- **Same values** → **same rank**
- **Skips numbers** after ties

Syntax

RANK() OVER (ORDER BY **column DESC)**

RANK() handles ties but leaves gaps

CASE in SQL

What is CASE?

- SQL's **if-else logic**
- Returns values based on conditions
- Works inside SELECT, WHERE, ORDER BY

Syntax

```
CASE
WHEN condition THEN
value
ELSE value
END
```

CASE brings decision-making into SQL

NoSQL Tutorial

Map-Reduce



SQL vs NoSQL

- SQL has tables with fixed rows and columns whereas NoSQL has dynamic data storage models (can be altered anytime).
- SQL is a general purpose database whereas NoSQL is a database used for handling documents, key-value pairs, wide-column data and graphs.
- SQL is fixed whereas NoSQL is flexible. (SQL has RDBMS schema consisting of relationships between tables and constraints defined initially which is not so for NoSQL - can be altered without handling relationships so any data can be handled with ease)



SQL vs NoSQL

- SQL has vertical scaling (increase resource of server where SQL is present - Central). NoSQL has horizontal scaling (data distributed among other local machines - Distributed).
- SQL has ACID properties. NoSQL has CAP properties (Consistency, Availability, Partition Tolerance)
- NoSQL has the power to handle very large amounts of data unlike SQL.



Examples

- SQL - MySQL, PostgreSQL, MS SQL Server
- NoSQL - MongoDB, Cassandra, HBase



Advantages of NoSQL


- **Flexible Databases:** Allows you to easily make changes to your database as requirements change.
- **Horizontal scaling :** Allow you to scale-out horizontally, meaning you can add cheaper commodity servers whenever you need to.
- **Fast Queries :** NoSQL does not require joins, unlike SQL, making execution of queries faster. Data in SQL databases is typically normalized, so queries for a single object or entity require you to join data from multiple tables. Increasing table size implies faster joins.

MapReduce

- MapReduce is a programming paradigm used for efficient processing in parallel over large data-sets in a distributed manner.
- The data is first split and then combined to produce the final result.

MapReduce - Working

The overall MapReduce word count process



MapReduce - Phases

- Its main use is to map the input data in key-value pairs.
- The input to the map may be a key-value pair where the key can be the id of some kind of address and value is the actual value that it keeps.
- The Map() function will be executed in its memory repository on each of these input key-value pairs and generates the intermediate key-value pair which works as input for the Reducer or Reduce() function.

MapReduce - Phases

- The intermediate key-value pairs that work as input for Reducer are shuffled and sort and send to the Reduce() function.
- Reducer aggregate or group the data based on its key-value pair as per the reducer algorithm written by the developer.

MapReduce - Working


- Input Phase – Here we have a Record Reader that translates each record in an input file and sends the parsed data to the mapper in the form of key-value pairs.
- Map – Map is a user-defined function, which takes a series of key-value pairs and processes each one of them to generate zero or more key-value pairs.
- Intermediate Keys – The key-value pairs generated by the mapper are known as intermediate keys.
- Shuffle and Sort – The Reducer task starts with the Shuffle and Sort step. It downloads the grouped key-value pairs onto the local machine, where the Reducer is running. The individual key-value pairs are sorted by key into a larger data list. The data list groups the equivalent keys together so that their values can be iterated easily in the Reducer task.

MapReduce - Working

- Reducer – The Reducer takes the grouped key-value paired data as input and runs a Reducer function on each one of them. Here, the data can be aggregated, filtered, and combined in a number of ways, and it requires a wide range of processing. Once the execution is over, it gives zero or more key-value pairs to the final step.
- Output Phase – In the output phase, we have an output formatter that translates the final key-value pairs from the Reducer function and writes them onto a file using a record writer.

MapReduce - Working


The overall MapReduce word count process



MapReduce - Concept of Combiner

- A combiner is a type of local Reducer that groups similar data from the map phase into identifiable sets. It takes the intermediate keys from the mapper as input and applies a user-defined code to aggregate the values in a small scope of one mapper. It is not a part of the main MapReduce algorithm; it is optional.

Combiner in Map - Reduce





EXAMPLE:

Find the total number of orders per country.

Sample Dataset

```
order_id, customer_id, country, category, price, quantity, order_date
1,1293, FRA, Clothing, 347.14, 3, 2020-09-17
2,370, GER, Clothing, 291.38, 4, 2015-06-18
3,807, FRA, Electronics, 118.83, 5, 2019-12-28
4,1938, USA, Sports, 133.48, 2, 2016-11-23
5,1090, BRA, Electronics, 350.16, 2, 2015-12-21
6,726, FRA, Books, 247.64, 2, 2017-05-07
7,1637, GER, Clothing, 459.33, 3, 2017-02-23
8,1759, FRA, Electronics, 87.06, 4, 2016-10-17
9,374, IND, Clothing, 275.97, 2, 2017-05-18
```



Mapper Code

```
import sys

filename = sys.argv[1]

with open(filename) as f:
    for line in f:
        if line.startswith("order_id"):
            continue
        fields = line.strip().split(",")
        country = fields[2]
        print(f"{country}\t1")
```

Combiner Code

```
import sys

current_key = None
current_count = 0

for line in sys.stdin:
    key, count = line.strip().split('\t')
    count = int(count)

    if current_key == key:
        current_count += count
    else:
        if current_key:
            print(f'{current_key}\t{current_count}')
        current_key = key
        current_count = count

if current_key:
    print(f'{current_key}\t{current_count}')
```



Reducer Code

```
import sys

current_key = None
current_count = 0

for line in sys.stdin:
    key, count = line.strip().split('\t')
    count = int(count)

    if current_key == key:
        current_count += count
    else:
        if current_key:
            print(f'{current_key}\t{current_count}')
        current_key = key
        current_count = count

if current_key:
    print(f'{current_key}\t{current_count}')
```



Bash File

```
#!/bin/bash

# Directory containing your data files
DATA_DIR="./data"

# Temporary file to store intermediate results
INTERMEDIATE_FILE="intermediate.txt"

# Output file for the final result
OUTPUT_FILE="result__.txt"

# Ensure the intermediate file is empty
> "$INTERMEDIATE_FILE"

# Loop through each text file in the data directory
for file in "$DATA_DIR"/*.txt; do
    # Process each file with the mapper, sort the output, then pass it to the combiner
    echo "Processing file: $file"
    python mapper.py "$file" | sort | python combiner.py >> "$INTERMEDIATE_FILE" &
done

# Wait for all background processes to complete
wait

# Sort the combined intermediate results and pass them to the reducer
sort -n "$INTERMEDIATE_FILE" | python reducer.py > "$OUTPUT_FILE"

# Clean up the intermediate file
rm "$INTERMEDIATE_FILE"
```



Result

BRA 9944

FRA 9994

GER 10075

IND 9968

USA 10019



THANK YOU